

Problem Set 3 Solutions

Department of Electrical and Computer Engineering
The University of Texas at Austin
EE 382N.1, Fall 2018
Instructor: Dam Sunwoo
TA: Pritesh Chhajed

Problem 1

Suppose we have the following loop executing on a pipelined LC-3b machine.

```
DOIT      STW    R1, R6, #0
          ADD    R6, R6, #2
          AND    R3, R1, R2
          BRz    EVEN
          ADD    R1, R1, #3
          ADD    R5, R5, #-1
          BRp    DOIT
EVEN      ADD    R1, R1, #1
          ADD    R7, R7, #-1
          BRp    DOIT
```

Assume that before the loop starts, the registers have the following **decimal** values stored in them:

Register	Value
R0	0
R1	0
R2	1
R3	0
R4	0
R5	5
R6	4000
R7	5

The fetch stage takes **one** cycle, the decode stage also takes **one** cycle, the execute stage takes a variable number of cycles depending on the type of instruction (see below), and the store stage takes **one** cycle.

All execution units (including the load/store unit) are fully pipelined and the following instructions that use these units take the indicated number of cycles:

Instruction	Number of Cycles
STW	3
ADD	3
AND	2
BR	1

Data forwarding is used wherever possible. Instructions that are dependent on the previous instructions can make use of the results produced right after the previous instruction finishes the execute stage. Multiple Instructions can write the results to the register file concurrently in the same cycle.

The target instruction after a branch can be fetched when the BR instruction is in ST stage. For example, the execution of an ADD instruction followed by a BR would look like:

ADD	F	D	E1	E2	E3	ST		
BR		F	D	-	-	E1	ST	
TARGET							F	D

The pipeline implements “in-order execution.” A scoreboarding scheme is used as discussed in class.

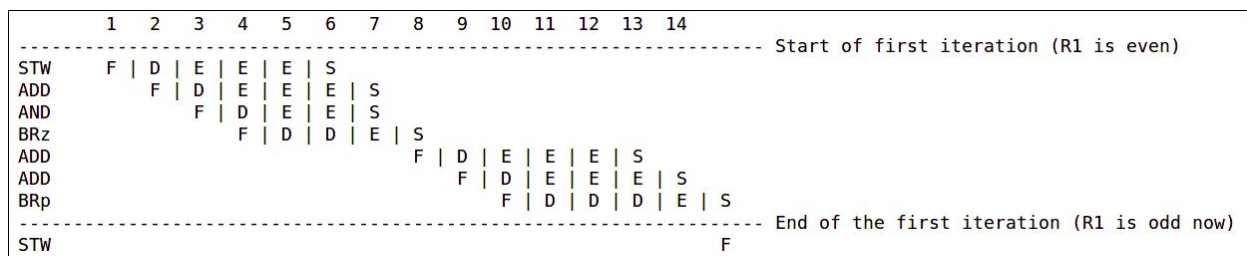
Answer the following questions:

1. How many cycles does the above loop take to execute if no branch prediction is used?
2. How many cycles does the above loop take to execute if all branches are predicted with 100% accuracy.
3. How many cycles does the above loop take to execute if a static BTFN (backward taken-forward not taken) branch prediction scheme is used to predict branch directions
4. What is the overall branch prediction accuracy? What is the prediction accuracy for each branch?

Assumptions:

- There are enough ports to the register file.
- There are enough ports to the memory.
- There are separate execution units for ADD, AND, STW, and BR instructions (They can all be in the execute stage at the same time.)

1. Assumption: Worst case scenario that branch target is always fetched during ST stage of BR instruction. If you assume that PC+2 is always fetched and in case of branch taken only, branch target fetched during ST stage is also acceptable.




The loop takes the same number of cycles to execute for even and odd values of R1. Each iteration takes 14 cycles in the steady state. There are 5 iterations for even values of R1 and 4 iterations for odd values of R1. The total number of cycles is:

$$(14 * 5) + (14 * 4) + 1 = 127$$

The extra 1 cycle comes from the last iteration (Store result stage of the BRp instruction).

2.




The loop takes the same number of cycles to execute for even and odd values of R1. Each iteration takes 13 cycles but 3 cycles can be overlapped with the next iteration. The total number of cycles is:

$$(10*9) + 3 = 93$$

3.

BRz instruction will always be predicted not taken. It is taken when R1 is even. So it will be mispredicted when R1 is even and correctly predicted when R1 is odd. The following diagram shows three consecutive iterations of the loop. In the first iteration, BRz is mispredicted, in the second iteration it is correctly predicted.

The first BRp instruction is always predicted taken. It is always predicted correctly. The second BRp instruction is also always predicted taken. It is mispredicted only once in the last iteration of the loop.



Loop steady state is shown above. It takes 22 cycles and it is repeated 4 times. The beginning of the loop (until the steady state) takes 10 cycles as shown above. The end of the loop (part of the last iteration which is not in steady state) takes 5 more cycles to execute. The total number of cycles is:

$$10 + (4 * 22) + 5 = 103$$

Prediction accuracies for each branch are:

- BRz = 4/9 (Correctly predicted when R1 is odd, R1 is odd for 4 iterations of the loop)
- first BRp = 4/4
- second BRp = 4/5 (Mispredicted only in the last iteration - Note that this misprediction does not affect the number of cycles it takes to execute the loop) v

Combined branch prediction accuracy = 12/18 = 67%

Problem 2

Consider the following program:

```

AND R0, R0, #0
BRz A
ADD R1, R0, #-1
C BRp B
A ADD R2, R0, #1
BRp C
B RET

```

A processor implements the GAg branch predictor as part of its microarchitecture. Assume the Branch History Register and Pattern History Table are as shown below before the program is run. The direction of the most recent branch is the right-most bit of the BHR; i.e., 1=taken, 0=not taken.

Branch History Register

101

Branch History Table

0	10
1	00
2	10
3	01
4	10
5	11
6	00
7	01

1. How many times does the branch predictor predict correctly?
2. What does the Branch History Register look like after the program finishes? The Branch History Table?

Assumption: Branch history register is updated speculatively (Based on Predicted value). While the Branch history table is updated based on actual value of branch.

The comments indicate whether the branch is predicted taken or not:

```
AND R0, R0, #0
BRz A           ; predicted taken, actually taken
ADD R1, R0, #-1
C BRp B         ; predicted not taken, actually taken
A ADD R2, R0, #1
  BRp C         ; predicted not taken, actually taken
B RET
```

Branch History Register

100

Branch History Table


0	10
1	00
2	10
3	10
4	10
5	11
6	01
7	01

Problem 3

Consider the following example used to explain Tomasulo's Algorithm:

```
Format: Opcode Destination Source1 Source2
MUL R3, R1, R2
ADD R5, R3, R4
ADD R7, R2, R6
ADD R10, R8, R9
MUL R11, R7, R10
ADD R5, R5, R11
MUL R10, R4, R10
```

Construct the Data Flow Graph for this program.



Problem 4

For the given assembly program, reverse-engineer the microarchitecture of the Tomasulo pipeline that would execute it as the execution timeline shown. Find the minimum cost solution (minimum number of reservation station entries, pipeline registers and etc).

Things that are known about the microarchitecture:

- There are one ADD unit, one MUL unit, and one DIV unit.
- The pipeline stages are IF, ID, (EXE/MEM), WB, where the latency of EXE/MEM stage varies among functional units.

j

Instr.	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
R1=R2*R3	IF	ID	M	M	M	M	M	M	M	WB													
R2=R1+R1		IF	ID	RS	RS	RS	RS	RS	RS	A	A	A	WB										
R4=R1+R2			IF	ID	RS	RS	RS	RS	RS	RS	RS	RS	A	A	A	WB							
R0=R0+R3				IF	ID	A	A	A	WB														
R3=R2+R3					IF	ID	ID	ID	ID	RS	RS	RS	RS	A	A	A	WB						
R5=R2*R5						IF	IF	IF	IF	ID	RS	RS	M	M	M	M	M	M	M	WB			
R1=R1/R0										IF	ID	D	D	D	D	D	D	D	D	D	D	WB	
R7=R5+R7											IF	ID	RS	RS	RS	RS	RS	RS	RS	A	A	A	WB
R6=R0/R2												IF	ID	RS	RS	RS	RS	RS	RS	RS	RS	D	D
R0=R2*R2													IF	ID	M	M	M	M	M	M	M	WB	

*A represents the ADD stage, M represents the MUL stage, and D represents the DIV stage.

a. What are the latencies of each functional unit?

ADD 3, MUL 7, DIV 10

b. What stages can be bypassed?

RS and WB

c. What would be the first instruction that would execute differently if your answer to the above was not the case?

$R1 = R2 * R3$ (1st instruction)

d. Are reservation stations centralized or distributed? Explain why the other can't be the case.

Distributed. The instruction ' $R3=R2+R3$ ' stalled the pipeline due to a structural hazard of the RS entry. For this to happen, the system should have either centralized RS with 2 entries or distributed RS with 2 entries for ADDER. Since 3 instructions are in the RSs at cycle 10, it should be the latter.

e. What would be the first instruction that would execute differently if your answer to the above was not the case? Assume the same number of RS entries total.

$R3 = R2 + R3$ (5th instruction)

f. How many reservation station entries are there? For centralized RS, show total number of entries. For distributed RS, show how many RS entries each FU has. Find the minimum cost solution.

ADD 2, MUL 1, DIV 1

Problem 5

Consider a fully-bypassed classic Tomasulo machine. There is an adder, a multiplier and a divider, each with its own single entry reservation station. All functional units are unpipelined. The register file has infinite read ports but only one write port and there is only one CDB (central data bus). There are Fetch, Decode, and Writeback stages. In the case of a conflict, the oldest instruction is always scheduled first.

Note: "fully-bypassed classic Tomasulo machine" means that an instruction can bypass the reservation station too and be scheduled on the functional unit if it is free.

Do the following:

1. Determine the latency of each of the functional units
2. Complete the timing diagram, indicating what stage/functional unit each instruction is in for each cycle (you must fill in every box between F and W in each row)

Functional Unit	Latency
Adder	4
Multiplier	1
Divider	2

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
R0 = R1 + R2	F	D	A	A	A	A	W							
R0 = R0 * R0		F	D	M ^R	M ^R	M ^R	M ^R	M ^R	M	W				
R1 = R1 / R2			F	D	/	/	/ ^W	W						
R2 = R3 * R2				F	D	M	M ^W	M ^W	W					
R3 = R1 / R3					F	D	/ ^R	/	/	/ ^W	W			
R0 = R0 / R2						F	D	/ ^R	/ ^R	/ ^R	/	/	W	
R0 = R2 + R2							F	D	A	A	A	A	A ^W	W

W Represents stalling due to CDB

R represents waiting in reservation station


Problem 6

A five instruction sequence executes according to Tomasulo's algorithm. Each instruction is of the form ADD DR,SR1,SR2 or MUL DR,SR1,SR2. ADDs are pipelined and take 9 cycles (F-D-E1-E2-E3-E4-E5-E6-WB). MULs are also pipelined and take 11 cycles (two extra execute stages). The microengine must wait until a result is in a register before it sources it (reads it as a source operand)

The register file before and after the sequence are shown below (tags for "After" are ignored).

Before				After			
	V	tag	value		V	tag	value
R0	1	z	4	R0	1		310
R1	1	z	5	R1	1		5
R2	1	z	6	R2	1		410
R3	1	z	7	R3	1		31
R4	1	z	8	R4	1		8
R5	1	z	9	R5	1		9
R6	1	z	10	R6	1		10
R7	1	z	11	R7	1		21


- Complete the five instruction sequence in program order in the space below. Note that we have helped you by giving you the opcode and two source operand addresses for instruction 4. (The program sequence is unique.)



2. In cycle 1 instruction 1 is fetched. In cycle 2, instruction 1 is decoded and instruction 2 is fetched. In cycle 3, instruction 1 starts execution, instruction 2 is decoded, and instruction 3 is fetched.

Assume the reservation stations are all initially empty. Put each instruction into the next available reservation station. For example, the first ADD goes into “a”. The first MUL goes into “x”. Instructions remain in the reservation stations until they are completed. Show the state of the reservation stations at the end of cycle 8.

Note: to make it easier for the grader, when allocating source registers to reservation stations, please always have the higher numbered register be assigned to SR2.



3. Show the state of the Register Alias Table (V, tag, Value) at the end of cycle 8.

	V	tag	value
R0			
R1			
R2			
R3			
R4			
R5			
R6			
R7			


Solution:

a.


1	ADD R7, R6, R7
2	ADD R3, R6, R7
3	MUL R0, R3, R6
4	MUL R2, R6, R6
5	ADD R2, R0, R2

b.

1	z	10	1	z	11	a
1	z	10	0	a	11	b
0	x	4	0	y	6	c



0	b	7	1	z	10	x
1	z	10	1	z	10	y
						z



c.

	V	tag	value
R0	0	x	4
R1	1	z	5
R2	0	c	6
R3	0	b	7
R4	1	z	8
R5	1	z	9
R6	1	z	10
R7	0	a	11

